
EXCISE: ONE-SHOT CAPABILITY EXTRACTION FROM OPEN LANGUAGE MODELS

PREPRINT

Arya Manjaramkar
arya@e3group.ai

June 2026

ABSTRACT

A deployed language model usually runs one capability while paying for all of them. We show that a single capability can be carved out of a frozen model and run on a small fraction of the network, as little as 1–2% of MLP channels for a well-covered task, in one short, label-free training run on a single consumer GPU. Learnable gates on every MLP channel train jointly with a low-rank adapter, held to the frozen base model by a forward-KL constraint, while a controller lowers the sparsity target as fast as behavior allows; the decision of when to stop is made by generation probes, not loss curves. The run needs no labels (the target is the model’s own output), reports a controller-found stopping point rather than requiring a fixed budget, and exports a physically smaller model: deleting the masked channels and pruning the vocabulary to the capability’s token support shrinks Qwen2.5-Math-1.5B from 1.54B to 0.17B parameters (9×) with no loss against masking. When a run exits at a configured lower bound or step cap, the reported floor is an upper bound on the behavioral-collapse minimum, not a claim that collapse has been reached. On 2-digit addition the run drives every seed to 1.2% of channels (0.7% given more steps) at over 91% verbatim reproduction, and recovers 100.9% of base accuracy at 5%. The same procedure, unchanged, also extracts function calling from Qwen3-4B (76% verbatim reproduction at 40% of channels), which we report as breadth rather than a head-to-head PRISM comparison. It exposes data as its dominant knob: regenerating a JSON-extraction substrate from 5× more diverse prompts drops its floor from 27.5% to 1.9% of channels at 97% fidelity. The single run replaces PRISM’s staged pipeline (adapt, attribute, select, re-train), whose extraction contract and recovery metric we adopt; on PRISM’s own arithmetic benchmark it passes the hand-tuned staged best on both axes (2.9% vs. 5.05% of channels, 97% vs. 91% recovery). Three negative results shaped the design, and we argue they generalize: distribution-level KL reads healthy while real recovery collapses, layer-granular gates let the adapter overfit silently, and a renormalized sparse KL quietly trains the unmasked model away from the base it is meant to anchor.

Keywords capability extraction · structured pruning · self-distillation · model compression

1 Introduction

Efficiency methods such as pruning, distillation, and quantization ask how much of a network can be removed while *aggregate* performance survives. A sharper question is this: given one named behavior, how small a part of the network can carry it once everything else is switched off? We answer it with a single training run that selects and relocates the behavior at the same time, needs no labels, reports a controller-found sparsity floor, and hands back a physically smaller model, all on one consumer GPU. The question itself, the *extraction contract* (a channel subset counts only if behavior survives with its complement deactivated), and *recovery* as the metric are due to PRISM [13], which we adopt and credit; its central finding is that a behavior-preserving adaptation step (“collimation”) can relocate a capability into far fewer channels than attribution alone suggests.

PRISM located that substrate with a staged pipeline: collimate, attribute, select a top- k mask, optionally re-collimate under the frozen mask. The stages cap what is reachable. An adapter trained *after* mask selection cannot shrink the mask; one trained *before* moves the frontier only indirectly; and the budget k is swept by hand, one costly round of

evaluations per point. Folding selection and relocation into a single objective removes both the ceiling and the manual sweep. PRISM itself names a training-time L_0 or gating objective as its highest-value open direction.

This paper presents that technique and ships it as a tool. Figure 1 shows how small the substrate gets. Our contributions:

1. **One-shot joint extraction.** Hard-concrete gates [11] on every MLP channel train jointly with a LoRA adapter [10], under a mass-covering KL constraint to the frozen base. Selection and relocation happen at once, and the split between “collimate-then-attribute” and “attribute-then-collimate” disappears (§2).
2. **Floor discovery by behavior probe.** A Lagrangian controller lowers the sparsity target whenever the KL budget is met; cheap greedy-generation probes veto the descent when the model’s actual outputs start to degrade. The full sparsity–fidelity frontier and its reported floor come from one run (§2).
3. **No labels.** The distillation target is the base model’s own greedy output and full distribution; no labeled data is required (§3).
4. **Physical export.** For gated MLPs, deleting a masked channel is exactly equivalent to zero-masking it, and pruning the vocabulary to the capability’s token support compounds the saving. We verify both inside every run and ship them: 1.54B $\rightarrow$ 0.17B parameters (9 $\times$) at undamaged fidelity (§3.4).
5. **Calibrated negative results.** Distribution-level KL is an unreliable stopping signal at high sparsity; a renormalized sparse-KL cache silently trains the unmasked model away from the base it should anchor; and layer-granular gates allow silent overfitting to the training distribution (§4). These reach beyond our method: any faithfulness metric built on truncated logit divergence inherits the same blind spots.

2 Method

Setup. Let M be a frozen decoder-only transformer and C its MLP intermediate channels ($L \times d_{\text{ff}}$, the inputs to each block’s down-projection), following PRISM’s unit of extraction. Each channel c gets a hard-concrete gate $z_c \in [0, 1]$ parameterized by $\log \alpha_c$ [11, 12]; a LoRA adapter θ (rank 32, all linear modules) supplies the capacity to relocate. Gates initialize from a $|\text{grad} \times \text{act}|$ attribution percentile rather than a constant. Attribution alone selects poor masks (Table 1), but it is a useful prior on the initial channel ordering. The teacher is the base model itself. Given prompts x that exercise the target capability, we take its greedy continuations $\hat{y} = M(x)$ and its full distribution $p_M(\cdot | x, \hat{y}_{<t})$ at each output position, cached once as top-128 sparse vectors *plus the residual off-support mass*. Keeping the residual is not bookkeeping. Renormalizing the support (the obvious implementation) makes the KL read $-\log(\text{coverage})$ even when student and teacher agree exactly, and puts positive gradient on every off-support token (§4).




Figure 1: Two-digit addition runs on 0.7–1.2% of MLP channels, found automatically in one label-free run. Qwen2.5-Math-1.5B. Red diamonds are **excise**’s controller stopping points: three configured lower-bound exits at 1.2% of channels and one extended step-cap exit at 0.71%, plotted by held-out verbatim self-match, a stricter metric that lower-bounds recovery. Gray squares are PRISM’s staged pipeline as reported [13] (recovery, task accuracy), shown for reference. Caveats in §5.




Figure 2: One-shot extraction. The frozen base model supplies label-free targets (greedy outputs and cached top- k distributions). Gates and adapter train jointly under a KL constraint, and a controller paced by generation probes finds the floor or records the configured exit. The run emits the sparsity–fidelity frontier, integrity receipts, and a physically sliced model.

Objective. Each step samples gates z (outside any gradient-checkpointed region; see §4) and minimizes

$$\mathcal{L} = \underbrace{\text{KL}(p_M \parallel p_{\theta,z})}_{\text{masked student tracks base}} + \beta_{ce} \text{CE}(\hat{y}) + \lambda \mathbb{E}[\|z\|_0] / |C| + \beta_g \underbrace{\text{KL}(p_M \parallel p_\theta)}_{\text{guardrail (every 4th step)}}, \quad (1)$$

with $\beta_{ce}=0.05$, $\beta_g=1$. Both KL terms are binned over the cached top- k support plus one residual bucket, which is zero exactly when the student matches the teacher. The forward (mass-covering) KL direction follows PRISM’s collimation analysis: the student may not drop probability anywhere the base places it. The guardrail evaluates the *unmasked* adapted model and anchors it to the base. This stops the adapter from “solving the task masked” by quietly becoming a different model; without it, unmasked accuracy drifted 3.7 points in our early runs. A train batch can only anchor points the adapter is already fitting, so the guardrail also alternates onto a rotating batch of off-task *anchor text* when provided, tunes its own cadence from the measured anchor KL, and stays active through the polish phase. Each of these closes a drift route we hit in practice (§3, §4).

Adaptive controller. A target open fraction τ starts at 1 and decays multiplicatively ($\times 0.993$ per step) whenever the EMA of the distillation KL sits under a budget (0.02–0.025 in all experiments); the Lagrange multiplier follows $\lambda \leftarrow \max(0, \lambda + \eta(\bar{z} - \tau))$. Once the open fraction drops below a threshold, every 150 steps a *probe* hardens the current top- k mask and checks whether a greedy rollout would still reproduce the teacher outputs. We compute this as teacher-forced argmax agreement (exactly equivalent by induction on the shared prefix, and $\sim 100\times$ cheaper than generating) on a dev split held out of the gradient batches, so the floor is decided on prompts the model has not memorized. The masked reading is compared against the unmasked model measured *at the same step*, never against an assumed 100%: batched bf16 decoding alone costs several points of verbatim self-match, and we must not blame masking for adapter drift. If the gap exceeds tolerance twice in a row, the floor is declared and the gates roll back to the last passing snapshot. If no probe failure occurs before a configured lower bound or step cap, we report that stopping point as a floor upper bound and record the exit reason in the receipts.

Once the descent stops, the mask is frozen at the floor budget and the adapter alone is polished briefly under the exact binary mask (guardrail still active), which removes the stochastic-gate train/eval mismatch. Final evaluation then hardens top- k masks at standard budgets. Because the polish specializes the adapter to the floor mask (we measured off-floor budgets reading *below* the floor without correction), each off-floor budget is briefly re-polished from the floor-polished snapshot before evaluation. From a single run this yields a frontier of deployable artifacts, not one point dressed up as many.

Receipts. Every run reports a random-mask control matched to the floor mask’s per-layer channel counts (which should read ≈ 0), unmasked drift, the base model’s own decode-noise self-match (the honest ceiling for every other number), the probe and guardrail traces, and binomial confidence intervals on held-out fidelity. Together they guard three failure modes (trivial tasks, cheating adapters, and masks that only correlate with behavior), two of which we hit and one we inherited from PRISM’s analysis.

3 Results

Every experiment ran on a single RTX 4090, through the released library’s public interface. Arithmetic gives the controlled comparison; JSON extraction and function calling test whether the same extraction machinery behaves on practical prompt-defined capabilities. This is the entire workflow:

```
from excise import extract, load_sliced

result = extract("Qwen/Qwen2.5-Math-1.5B", prompts=prompts) # no labels
print(result.report()) # frontier, floor, receipts, CIs
small = result.export_sliced(prune_vocabulary=True) # 1.54B -> 0.17B
result.save("substrate/") # receipts + the sliced artifact
model = load_sliced("substrate/sliced") # round-trips
```

Code, configs, and per-run receipts: github.com/Aryagm/excise.

3.1 Arithmetic: extracting addition into around 1% of the network

We use PRISM’s testbed: 2-digit addition on Qwen2.5-Math-1.5B [16] (canonical prompt "a + b =", exhaustive sums, 10% held out, 250,880 MLP channels), zero-isolation deactivation, and recovery = masked accuracy / base accuracy (base: 97.4%). Table 1 gives the key points, and Figure 1 places the library’s controller stopping points against the staged pipeline. Table 1 reports task-accuracy recovery frontier points; the lower controller exits in Figure 1 report held-out verbatim self-match, a stricter metric.

Table 1: Arithmetic extraction. Every seed beats the staged pipeline at 5%, and the lowest $\geq 95\%$ recovery point is at or below the hand-refined minimal circuit with higher recovery. The unmasked guardrail held base accuracy exactly (97.4%) in all runs.

	channels kept	recovery
PRISM raw attribution mask	5.75%	29.0%
PRISM staged collimation	5.05%	91.3%
PRISM refined minimal circuit (manual)	4.65%	90.6%
random mask control (ours)	5%	0.0%
grad×act attribution mask (ours)	5%	0.3%
excise, one run (seed 42 / 43 / 44)	5%	100.9 / 100.4 / 99.5%
excise, lowest $\geq 95\%$ recovery point	2.9 / 3.8 / 4.1%	97.2 / 95.4 / 97.9%

Two ablations (Figure 3): restricting the adapter to MLP projections (the strictest frozen-scaffold contract) costs ~ 3 points at 5% and runs 25% faster; removing the hard-token cross-entropy term ($\beta_{ce}=0$) still reaches 96.7% recovery but floors shallower (5.3%). Hard-token CE is not needed for validity, only for the last $\sim 2\%$ of budget.

The reported floor is set by compute, not by behavior. All three seeds descend to the controller’s configured lower bound of 1.2% of channels, with the masked–unmasked probe gap inside tolerance throughout, at 91.0–92.0% held-out verbatim self-match. Lowering the bound and giving one seed 9,000 steps reaches 0.71% of channels ($\sim 1,800$ of 250,880) at 91.1%, exiting on the step cap while still descending at ~ 0.1 points per thousand steps (Figure 1). When the prompts cover the task exhaustively, these exits are controller-found stopping points, not measured behavioral-collapse minima: the minimal circuit for 2-digit addition is somewhere below 0.7%, $6.5\times$ under the published hand-refined one.

3.2 Function calling: a real-world capability, label-free

We extract single-turn function calling from Qwen3-4B [17] using BFCL prompts [14] (400 tasks; 391 yield clean tool calls from the base model; 320 train, 80 held out; 350,208 channels). No gold calls are used: the teacher is the model’s own greedy tool call, and fidelity is *verbatim* self-match, stricter than BFCL’s semantic scoring, since one diverging token fails the example. Held-out fidelity reaches 77.5% at 50% of channels and 76.3% at 40% (Figure 4). Two things stand out.

First, the *unmasked* adapted model self-matches only 86% on held-out prompts. Masking is not the main source of loss; adapter drift under a 320-prompt training set is. Against that ceiling, the masked model keeps $\sim 90\%$ at 40–50% budgets.




Figure 3: The result is seed-stable and survives ablation. Recovery at 5% of channels across seeds and ablations; channel floors found by the controller annotated within bars.




Figure 4: Function calling extracts label-free; data, not machinery, binds generalization. Qwen3-4B. Left: probe trace during descent; exact self-match holds at 92–98% from 61% down to 12% open on training prompts. Right: held-out self-match by budget; the dashed line is the unmasked adapted model’s own self-match (86%), which upper-bounds every masked number.

Second, probe fidelity on training prompts stays at 92–98% down to 12% open, far above held-out fidelity. That gap is a generalization effect an exhaustively-covered task like arithmetic cannot show: the binding constraint is prompt diversity, not the extraction machinery. To test this, a second experiment scales the prompts 6× and hardens the task (all six single-turn BFCL splits, plus 1–2 distractor tools per prompt, so the substrate must also carry tool *selection*). Drift improves to 88.4% despite the harder task, and the binding constraint moves again: dev-split probes stop the descent at 60.7% of channels (75.3% held) after only five epochs. Given enough data, function calling is bound by training steps, not by memorization. The two settings solve different tasks (one given tool vs. selection among distractors), so we do not compare them directly. For calibration, PRISM’s raw mask recovered 19.1% at 36.2% of channels on the 8B sibling model, and its best staged collimator (with a curated, schema-validated gold corpus and a multi-run objective search) reached 84.6%; metric and model-size differences make this directional context, not a head-to-head (§5).

3.3 Breadth: other architectures and capabilities

To check that the method is a tool and not a Qwen-arithmetic demo, we ran the *released library* (not the research scripts) on two further settings.

SmolLM2-1.7B (Llama architecture) [2]. Few-shot 2-digit addition extracts to a 7.4% floor at 97.2% held-out verbatim self-match, sliced 1.75B $\rightarrow$ 0.59B (3.0×), all receipts clean.

JSON extraction (Qwen2.5-1.5B-Instruct, chat-formatted prompts). From 3,000 label-free prompts spanning eight sentence forms and four instruction phrasings, structured JSON extraction floors at 1.9% of channels at 97.2%




Figure 5: Same capability, same model, same config; only the prompts change. JSON extraction from 600 single-template prompts vs. 3,000 varied ones (eight sentence forms, four instruction phrasings): the floor collapses 14 $\times$, fidelity at the floor rises seven points, and unmasked drift nearly disappears. Both runs are label-free, so the diverse variant costs nothing but generation.




Figure 6: Where the parameters go. Qwen2.5-Math-1.5B before and after extraction at the 0.71% floor: 99.3% of MLP channels and 98.9% of the vocabulary are deleted outright; attention and norms, which the method does not touch, are 94% of what remains, the obvious next structured dimension to attack.

held-out self-match (still step-capped), with 1.7 points of unmasked drift, a flat $\sim 98\%$ frontier at every budget, and a verified $1.58\text{B} \rightarrow 0.21\text{B}$ export. The capability touches 9,448 token ids; broader capabilities are broader in vocabulary too.

Data coverage is the dominant knob. Restricting the same capability to 600 prompts from a *single* sentence template inflates the floor 14 $\times$ (to 27.5%) and drops fidelity to 90% (Figure 5). Since the held-out prompts share that template, even the 90% partly measures template fit. We first read our settings as “capability breadth tracks substrate size”; this ablation corrects that. Substrate size measures the capability *as the prompts exercise it*. Floors are upper bounds set jointly by data coverage and compression compute, not intrinsic capability sizes, and data is the cheapest lever, since label-free prompts cost nothing to diversify.

3.4 Physical export: slicing is masking, and the vocabulary is next

For gated MLPs, a zero-masked channel contributes nothing to the down-projection, and a deleted one contributes nothing by absence. The two are mathematically identical, so the masked evaluations above transfer to a physically smaller model. We verify it: slicing Qwen2.5-Math-1.5B at the 2.9% floor gives 94.8% accuracy against 94.7% masked (one example in 810), at $1.54\text{B} \rightarrow 0.42\text{B}$ parameters (3.66 $\times$).

After MLP slicing, the vocabulary is the dominant remaining cost. 2-digit addition touches only 1,746 of 151,936 token ids, so slicing the embedding rows to that support takes the same model to 0.17B parameters (9.0 $\times$) at the 0.71% floor (Figure 6); the exported model speaks a remapped id space, with the mapping stored in its config. Both equivalences are re-verified inside every run and logged in the receipts (masked, sliced, and vocabulary-pruned self-match agreed to within one evaluation example here), and the library round-trips the heterogeneous-width artifact (`excise.load_sliced`).

Slicing buys memory, not speed at this scale, and we say so plainly because such claims are routinely overstated. Our original benchmark showed only a 1.11 $\times$ generation speedup, and a careful decode-throughput benchmark (batch

1–64, 64-token greedy decode, eager) shows no speedup at all: the 0.17B and 1.54B models both decode at ~ 40 tokens/s/sequence on an RTX 4090, because single-stream decoding there is bound by launch overhead, not compute. Memory shrinks unconditionally; latency gains would need larger models, batched serving, or a compiled runtime.

4 What fails, and what it implies

The method of §2 is an end state. Most of its load-bearing choices exist because an earlier, more obvious implementation failed in a way worth recording, and this section is that record. Each failure was caught by the receipts the tool itself emits, and each says something beyond this method.

Distribution-level KL is miscalibrated at high sparsity. Our first controller paced the descent on the KL budget alone. It reached 2.2% open with EMA KL comfortably under budget, while true recovery collapsed to 62–66% (Figure 8). Teacher-forced divergence on answer tokens simply does not capture the way error compounds across an autoregressive rollout. The fix, periodically generating and exact-matching, costs seconds per probe and is the load-bearing stability feature of the method. The lesson extends past extraction: any faithfulness or circuit-quality metric built on teacher-forced logit divergence inherits the same blind spot.

Renormalized sparse KL quietly sharpens the model it should anchor. Caching the teacher’s top- k probabilities *renormalized* (the obvious implementation, and our first) biases every loss built on it. With the student exactly equal to the teacher, the “KL” reads $-\log(\text{coverage})$ instead of zero, and its gradient $s_j - \hat{p}_j$ is positive on all off-support mass, pushing the student to concentrate onto each batch’s cached support. On peaked tasks (arithmetic, coverage ≈ 1) the bias vanishes, which is how it survived three validated batteries. On diffuse tasks it surfaces in the worst place: the guardrail, whose entire job is to hold the unmasked model *at* the base, instead trains it to sharpen. The receipts caught it: a JSON run whose unmasked self-match read 55.8% (made worse by a then-guardrail-free polish phase). With the residual bucket restored, anchor batches added, and the guardrail active through polish, the same task reads 88.3% against a 96.7% decode-noise ceiling, and its floor *improves*. Figure 7a shows why train-batch monitoring missed it: at the same step, the guardrail KL reads 3×10^{-4} on train batches (apparently perfect) and 5×10^{-3} –0.15 on off-task anchors. Drift lives exactly where a train-only guardrail cannot see it; an anchor that watches only the training distribution anchors nothing. The lesson mirrors the one above: a truncated divergence is not a divergence, and it fails precisely where its value is supposed to be zero.

Probes calibrated on memorized data flatter the floor. Our first release probed on prompts that were also in the gradient batches, against an assumed-perfect baseline. But even the unmodified base model reproduces its own batched-bf16 greedy targets only $\sim 92\%$ verbatim, so part of every probe reading was decode noise, not masking damage (Figure 7b gives the corrected decision rule). Re-running the identical extraction with dev-split probes, scored against the unmasked model at the same step, with rollback to the last passing sparsity and the binned KL, moves every number the tool reports (Table 2); most of this paper’s headline results were out of reach under the original calibration. The lesson is quantitative: with $n=64$ verbatim probes, binomial noise alone is ± 4 points, so a 3-point tolerance sits inside the noise band. The first two probes of a run tripped it and froze the floor at 7.9%, where a noise-aware tolerance descends to 1.2% with probes passing throughout. Tolerances on generation metrics must be set against the metric’s measured noise floor, not chosen for strictness.

Table 2: The same extraction under miscalibrated vs. calibrated receipts (arithmetic; identical task, seed, hardware, and metric; base decode-noise ceiling 99.9%). Calibration is not cosmetic: it halves measured drift and reaches a $6\times$ lower floor.

	original	calibrated (3 seeds)
floor found (probe-governed)	7.6%	1.2 / 1.2 / 1.2%
held-out self-match at floor	89.0%	91.4 / 92.0 / 91.0%
unmasked drift from base	10.7 pts	6.2 / 5.8 / 7.0 pts
exported parameters (from 1.54B)	0.48B	0.17B

Layer-granular gates permit silent overfitting. Adding per-layer gates (hierarchical L_0 , meant to enable wall-clock-friendly whole-layer drops) closed 16 of 28 layers. Training KL stayed low because the adapter rerouted around the missing layers *on the training distribution*, and held-out recovery collapsed to 66%. Channel-granular gates with all layers present do not show this. Structured sparsity for speed should be grafted on after extraction, not learned jointly.




Figure 7: Two failure modes, instrumented (JSON extraction run). (a) Guardrail KL on train batches (gray) vs. rotating off-task anchor batches (red), same adapter and same steps: the train-batch view reads 10–100× lower than the anchor view. Drift lives exactly where a train-only guardrail cannot see it. (b) Floor detection: the masked probe (red) is compared against the unmasked model measured at the same step (gray, itself depressed by decode noise and residual drift), never against an assumed 100%. When the masked probe falls below tolerance twice, the gates roll back to the last passing sparsity.




Figure 8: Losses lie; generations don't. An early KL-only controller: EMA distillation KL stays under budget throughout the descent (left axis) while true held-out recovery collapses (crosses, right axis). Generation probes detect this; losses do not.

Engineering notes. Stochastic gate sampling must happen outside gradient-checkpointed regions: in-region RNG produces a different recomputation graph and breaks the backward pass. Pre-sampling once per step costs nothing and composes with checkpointing, which 4B+ models need on consumer GPUs. Separately, `grad×act` attribution must accumulate under `no_grad`: an in-place `+=` of a grad-requiring product turns the accumulator into an autograd node that retains every iteration’s activation graph. That is megabytes and invisible on 10-token arithmetic (it survived every validated battery), but gigabytes per iteration and an instant OOM on 100-token JSON. Bugs that scale with sequence length hide in short benchmarks.

5 Honest comparison notes

- **Baselines.** Our arithmetic baseline reproduces PRISM’s qualitative finding (attribution masks fail small: 0.3% recovery at 5%) but uses `grad×act` attribution and zero-isolation, not their ReLP and counterfactual patching. PRISM’s reported numbers come from their paper, not our re-runs, and zero-isolation is the harsher operator, so our baseline understates theirs.
- **Function calling.** The comparison differs in model size (4B vs. 8B) and metric (verbatim self-match vs. BFCL exact match); we present it as context, not a head-to-head.
- **Seeds.** Arithmetic results are 3 seeds. Function calling is a single seed, and its floor was set by the step cap, not detected by a probe.

- **Scope.** PRISM’s 10^3 H100-hour budget bought breadth (translation, an 8B model, extensive controls) that ours, at under \$10 per result, did not. The claim here is that the *per-result* cost collapses, not that our evidence is broader.

6 Related work

PRISM [13] defines the problem, contract, and metric we use, and our method is its stated future direction, built. Learnable L_0 masks during adaptation descend from Louizos et al. [11] and movement pruning [15]; differentiable circuit discovery applies the same machinery to interpretability [3, 4]. What sets us apart from both is the objective: a behavior-preservation contract with an unmasked-validity guardrail, rather than aggregate-loss compression or circuit faithfulness. Self-distillation into a substructure relates to on-policy and generalized knowledge distillation [1, 9], here with teacher and student sharing weights and the student differing only by a mask. Superposition [6] explains why attribution-only masks fail and adaptation is needed; sparse autoencoders [5] and weight-sparse training [8] change the basis instead. The lottery-ticket line [7] seeks subnetworks that *retrain* to full performance, whereas extraction requires the behavior to survive *without* any retraining from initialization.

7 Limitations

Capabilities are bounded by prompt coverage: the tool can flag a generalization gap (through held-out receipts) but cannot close it without more data. Self-match undercounts semantic fidelity. Vocabulary-pruned artifacts speak a remapped id space, and the function-calling vocabulary support is 101k of 152k ids, so that lever pays little there. Attention is untouched and dominates the exported model. We have not yet run the 8B headline, multi-seed function calling, or non-Qwen replications beyond SmoLLM2 and small architecture tests; these gate any stronger claims. Easier extraction also eases misuse, such as pulling narrow capabilities out of safety-tuned models, so receipts and intended-use documentation ship with the tool, following PRISM’s stance.

8 Conclusion

Capability extraction can be a single, cheap, label-free training run. It selects and relocates a behavior at once, finds or bounds its floor by generating rather than by watching losses, and leaves behind a deployable frontier, integrity receipts, and a model $9\times$ smaller, all on one consumer GPU. We build on PRISM’s contract (behavior must survive with the complement switched off) and pass its staged pipeline on the arithmetic benchmark, but that comparison is not the point. The durable lessons are that a capability is far smaller and cheaper to isolate than a staged search suggests, and that logit-level divergence is not behavior: any metric built on it will miss the collapses a cheap generation probe catches. Extraction is now a `pip install`, not a pipeline.

Reproducibility

All numbers in this paper were produced by the released library or the research scripts in the same repository, on one rented RTX 4090. The paper’s results correspond to commit `b0eb63aabe3c` / release `v0.2.0`. The paper’s figures regenerate from committed result JSONs via `paper/make_figures.py`.

References

- [1] Rishabh Agarwal, Nino Vieillard, Yongchao Zhou, Piotr Stanczyk, Sabela Ramos Garea, Matthieu Geist, and Olivier Bachem. On-policy distillation of language models: Learning from self-generated mistakes. In *ICLR*, 2024.
- [2] Loubna Ben Allal et al. SmoLLM2: When smol goes big – data-centric training of a small language model. *arXiv:2502.02737*, 2025.
- [3] Adithya Bhaskar, Alexander Wettig, Dan Friedman, and Danqi Chen. Finding transformer circuits with edge pruning. In *NeurIPS*, 2024.
- [4] Arthur Conmy, Augustine N. Mavor-Parker, Aengus Lynch, Stefan Heimersheim, and Adrià Garriga-Alonso. Towards automated circuit discovery for mechanistic interpretability. In *NeurIPS*, 2023.
- [5] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv:2309.08600*, 2023.

- [6] Nelson Elhage et al. Toy models of superposition. *Transformer Circuits Thread*, 2022.
- [7] Jonathan Frankle and Michael Carbin. The lottery ticket hypothesis: Finding sparse, trainable neural networks. In *ICLR*, 2019.
- [8] Leo Gao et al. Weight-sparse transformers have interpretable circuits. *arXiv:2511.13653*, 2025.
- [9] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv:1503.02531*, 2015.
- [10] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. Lora: Low-rank adaptation of large language models. In *ICLR*, 2022.
- [11] Christos Louizos, Max Welling, and Diederik P. Kingma. Learning sparse neural networks through l_0 regularization. In *ICLR*, 2018.
- [12] Chris J. Maddison, Andriy Mnih, and Yee Whye Teh. The concrete distribution: A continuous relaxation of discrete random variables. In *ICLR*, 2017.
- [13] Abhishek Mishra and Krishna Pagare. Prism: Unlocking language model capability extraction. In *1st Conference For AI Scientists (CAISc)*, 2026.
- [14] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The berkeley function calling leaderboard (bfcl): From tool use to agentic evaluation of large language models. In *ICML*, 2025.
- [15] Victor Sanh, Thomas Wolf, and Alexander M. Rush. Movement pruning: Adaptive sparsity by fine-tuning. In *NeurIPS*, 2020.
- [16] An Yang et al. Qwen2.5-math technical report: Toward mathematical expert model via self-improvement. *arXiv:2409.12122*, 2024.
- [17] An Yang et al. Qwen3 technical report. *arXiv:2505.09388*, 2025.